



INTRODUCTION TO ARTIFICIAL INTELLIGENCE  
SWE3011\_41  
FALL 2023

---

Project  
Anonymous ACL submission

---

Sungkyunkwan University

November 30, 2023

---

## Abstract

Within the context of this assignment, a substantial hands-on encounter with applications of artificial intelligence will be attained, with a particular focus on two pivotal tasks: traditional supervised learning methodologies and prompt engineering. Upon the culmination of this assignment, an adept understanding of the practical implementation of artificial intelligence techniques for the purpose of classification shall be cultivated. Furthermore, through collaborative participation in the curation of prompts and the execution of experiments, the acquisition of insights and skills that are notably advantageous for a group project is anticipated.

## 1 Introduction

In this lab journal, the primary objective is to offer a comprehensive practical experience in the field of artificial intelligence, with a specific focus on supervised learning. Upon the conclusion of this project, a solid and pragmatic understanding of traditional machine learning techniques for text classification is expected to be acquired. Furthermore, the utilization of Langchain for prompt engineering tasks, especially in the context of Large Language Models (LLMs), will have been explored. This multifaceted approach will equip us with a versatile skill set and a deeper insight into the practical aspects of artificial intelligence.

This assignment categorizes tasks into three main problems:

1. **Supervised Text Classification using Traditional Machine Learning Methods:** The task involves implementing a text classification process using traditional machine learning methods. A dataset and skeleton code utilizing the scikit-learn library, with specific sections left blank for completion, will be provided. The responsibility is to finalize the code, conduct experiments, and evaluate to derive results.

2. **Prompt Engineering for LLMs via Langchain Framework:** Using Langchain, a tool designed for crafting prompts for language models, a series of prompts will be created to elicit specific and varied responses. Baseline code for conducting prompt engineering experiments with ChatGPT and Large Language Models on Hugging Face Hub is provided. The task involves completing the code and optimizing the prompts for sentiment classification.

3. **LATEX:** The third focus of this assignment is report writing using LATEX, a document preparation system that aids in creating structured documents. This part helps students develop skills in professional communication and documentation within the AI field.

This journal serves as a documentation of the academic journey in addressing these AI challenges and acquiring practical skills in the field of artificial intelligence.

## 2 Dataset

**The GLUE benchmark** stands as a compendium of diverse Natural Language Processing (NLP) tasks, employed to assess the efficacy of language models. A significant component of the GLUE benchmark, known as Stanford Sentiment Treebank V2 (SST-2), is dedicated to evaluating the performance of models in sentiment classification on the SST-2 dataset and other related tasks, serving as a litmus test for their broader language comprehension and reasoning capabilities (Wang et al., 2018).

**The SST-2 dataset** is a renowned benchmark dataset in the realm of NLP, frequently employed for sentiment analysis undertakings. An extension of the Stanford Sentiment Treebank (SST), which originally featured fine-grained sentiment labels encompassing categories such as very positive, positive, neutral, negative, and very negative. SST-2 consists of sentences sourced from movie

---

reviews, each meticulously annotated with binary labels signifying either positive or negative sentiment (Socher et al., 2013).

### 3 Methodology

For both Task 1 and Task 2, distinct methods have been chosen to address their specific objectives.

For Task 1, which centers on supervised text classification, the chosen trio of methods comprises logistic regression, random forest, and naive Bayes classifiers. These algorithms were meticulously selected for their suitability and proven effectiveness in tackling text classification challenges.

For Task 2, which focuses on prompt engineering for Large Language Models, LangChain was chosen as the tool to design both zero-shot and few-shot prompts. These prompts were to be tested using ChatGPT and the Large Language Models available on the Hugging Face Hub in order to solve the sentiment classification problem on the SST-2 dataset.

It's important to note that, for this task, the decision was made to abstain from using OpenAI models. Instead, the focus was solely on experimenting with freely available models on the Hugging Face Hub. This approach allowed exploration of the capabilities of large language models within the constraints of free resources.

## 4 Experiments

### 4.0.1 Task 1

In the experimental phase of this project, a comprehensive analysis was undertaken to evaluate the effectiveness of three distinct machine learning models: Logistic Regression, Random Forest, and the Naive Bayes Classifier. The approach adopted was two-fold, entailing not only model selection but also the intricate process of hyperparameter tuning and cross-validation. In this context, the Logistic Regression model demonstrated remarkable performance, optimizing its hyper-

parameters to yield the most accurate results. Specifically, the observation revealed that a regularization strength parameter, denoted as 'C,' was fine-tuned to the value of 100, and the regularization term 'penalty' was set to 'l2'. In essence, 'C' was controlled to determine the degree of regularization, with larger values signifying less regularization and a potentially higher sensitivity to the training data, while 'penalty' indicated the type of regularization applied, with 'l2' corresponding to the L2 penalty. On the other hand, when the Random Forest model was employed, the hyperparameter optimization revealed that the ideal configuration consisted of a maximum depth set to 'None' and an ensemble size, or 'n\_estimators,' fixed at 200. Here, 'max\_depth' signified the depth of each tree in the Random Forest, with 'None' indicating no restrictions on the depth. As for 'n\_estimators,' it represented the number of trees within the ensemble. Lastly, in the case of the Naive Bayes Classifier, it was determined that the hyperparameter 'alpha' was most effective at a value of 0.1, indicating the Laplace smoothing factor for the model. Through conducting this detailed experimentation and elucidating the significance of these parameter choices, the objective was not only to optimize model performance but also to gain a profound understanding of the underlying mechanisms that govern these machine learning algorithms.

### 4.0.2 Task 2

For the experimental part of Task 2, both zero-shot and few-shot prompts were prepared in order to utilize the Hugging Face Hub's Large Language Models' ability to perform sentiment analysis of the SST-2 dataset. "Zero-shot prompt" refers to prompt engineering where the model is not trained with examples on how to perform the given task, relying only on its existing knowledge, while "few-shot prompt" refers to prompts where the model is trained with specific examples on how to complete the task. For both the zero-shot and few-shot prompts, the model was

---

given the following prefix, as the instruction for its task of sentiment classification:

*"You are a sentiment classifier. Classify the given text input as "positive" or "negative" based on its tone. Please read each sentence carefully and determine its sentiment based on common definitions.*

- *<positive>: A positive sentiment indicates a favorable or optimistic tone. It expresses approval, satisfaction, or positive emotions.*
- *<negative>: A negative sentiment indicates an unfavorable or pessimistic tone. It expresses disapproval, dissatisfaction, or negative emotions.*

*Please read each sentence and classify it as either "positive" or "negative" based on the provided definitions."*

Apart from the prefix that serves as the task description, for the few-shot prompt, the model was trained with a provided dataset that shares similar characteristics to the SST-2 dataset. Due to the constraints of using only free resources, it was not possible to provide the Large Language Model with more than 27 training examples per execution. Since the training dataset was larger than that, a different set of examples was used for each execution. Each of the examples in the set was selected randomly out of all the available ones in the dataset; however, in order to maximize the accuracy of the Large Language Model predictions, a constraint was applied to maintain consistent proportions of positive and negative instances in each set. This approach was taken to prevent any imbalance between positive and negative examples in the generated sets, which could impact the model's training. After executing the prompts multiple times, it became apparent that the accuracy of few-shot prompts varied with each iteration due to changing examples influencing the model. Consequently, three iterations were

conducted to capture the accuracy range, constrained by limited resources. In the case of zero-shot prompts, a single iteration was performed for assessment. The resulting accuracy of the zero-shot prompt and the accuracy ranges of the few-shot prompts were compared in chapters "Results" and "Discussions".



---

## 6 Discussions

### 6.0.1 Task 1

Accuracy was chosen as the primary evaluation metric for Task 1, aligning with the approach taken in the GLUE benchmark. To evaluate the models, the provided test\_dataset was employed.

It is evident that the logistic regression model outperformed the other two models, namely random forest and naive Bayes, with the highest accuracy of 0.77. The logistic regression model exhibited a balanced performance in terms of precision, recall, and F1-scores for both classes 0 and 1. This indicates its effectiveness in handling text classification tasks. On the other hand, the random forest model achieved an accuracy of 0.71. While it provided competitive results, its performance was slightly lower compared to logistic regression. The classification report demonstrated reasonable precision, recall, and F1-scores for both classes, making it a viable choice but not the top-performing one.

Lastly, the naive Bayes model exhibited an accuracy of 0.70, which was consistent with random forest but slightly lower than logistic regression. The classification report also showed balanced precision, recall, and F1-scores for both classes.

### 6.0.2 Task 2

Accuracy was also chosen as the primary evaluation metric for Task 2, in order to stay consistent with the findings of Task 1, as well as provide insight into the capability of the Large Language Model to solve the sentiment classification problem on the SST-2 dataset.

The most important result was the difference in accuracy when employing few-shot prompts instead of zero-shot prompts. This was consistent with our expectation that sentiment classification yields better results when the model is trained with specific examples rather than when the task is simply performed based on the model’s pre-existing knowledge. This poses the idea that few-shot prompts are superior for solving sentiment classification problems. Because of this, choosing a proper

dataset and improving the quality, diversity and proportionality of positive and negative emotion examples is an important part of solving problems like this. Lastly, while there is a positive correlation between the amount of examples provided and the accuracy of the predictions, it is important to not overtune the model with an excessive amount of examples.

## 7 Conclusion

### 7.0.1 Task 1

In conclusion, this project delved into the application of various machine learning models, including logistic regression, random forest, and naive Bayes, for the task of text classification. Through rigorous experimentation and evaluation, it was observed that the logistic regression model emerged as the top-performing model, achieving the highest accuracy and demonstrating balanced precision, recall, and F1-scores for both classes. This emphasizes the significance of selecting an appropriate machine learning model tailored to the task at hand.

Despite the success of the logistic regression model, there are potential areas for improvement. First and foremost, further exploration of feature engineering techniques and text preprocessing methods could enhance the performance of all models. Additionally, the incorporation of more advanced machine learning models or deep learning architectures may lead to improved results in text classification tasks. Furthermore, the dataset size and diversity could be increased to bolster the models’ ability to generalize across different textual data.

Moreover, an investigation into hyperparameter tuning techniques, such as grid search or Bayesian optimization, could be beneficial to fine-tune the models further. Finally, ensembling methods, which combine the strengths of multiple models, may provide an opportunity to boost overall performance.

In summary, while logistic regression proved to be the most effective model in this specific context, there remains room for enhancements through various avenues, in-

---

cluding feature engineering, model selection, and hyperparameter tuning. This project serves as a valuable foundation for future research and applications in the field of text classification and natural language processing.

### 7.0.2 Task 2

In conclusion, Task 2 of this project allowed us to explore prompt engineering and how it is a fundamental part of AI projects. In this case, a sentiment classification problem was solved through zero-shot and few-shot prompts. It was undeniable that the few-shot prompt performed significantly better than the zero-shot prompt for this task. Not only that, but is also relevant to point out that the few-shot prompt achieved a range of accuracy at sentiment classification superior to that achieved by the machine learning methods utilized in Task 1. This could be attributed to many factors, most notably the nature of the SST-2 dataset. This dataset is text based, which makes it well suited to the use of prompt engineering. The machine learning methods explored in Task 1 are better suited for dataframes with attributes, especially Logistic Regression, which had the biggest accuracy rate in Task 1 and works better with numerical datasets.

## References

- Momojit Biswas. 2020. Prompt engineering: Unleashing the power of few-shot learning for accurate spam detection. Medium.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1631–1642, Seattle, Washington, USA. Association for Computational Linguistics.
- Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. 2018. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 353–355, Brussels, Belgium. Association for Computational Linguistics.

---

## Code

### Task\_1

```
1  # -*- coding: utf-8 -*-
2  """SWE3011_41_Task1.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1RKXBhSWJVU9MqRcyZBCIKWBRx6SMf_d
8
9  # SWE3011_41 Task1
10
11  **Supervised Text Classification using traditional machine learning methods**
12
13  1. Complete all the functions given.
14  2. Conduct various experiments including hyper-parameter tuning, cross validation, etc.
15  3. Write a report on the analysis of experiment results.
16
17  **0. Installation**
18
19  **1. Load Dataset**
20  """
21
22  pip install datasets
23
24  """
25  Evaluation should be done using **provided test dataset**"""
26
27  from datasets import load_dataset
28
29  train_ds = load_dataset("glue", "sst2", split="train")
30
31  # Evaluation should be done using test_ds
32  test_ds = load_dataset("csv", data_files="./test_dataset.csv")['train']
33
34  """**2. Preparing Dataset**"""
35
36  from sklearn.feature_extraction.text import TfidfVectorizer
37
38  def transform_data(X_train, X_test):
39      """
40      Input:
41      - X_train, X_test: Series containing the text data for training and testing respectively.
42
43      Output:
44      - X_train_tfidf, X_test_tfidf: Transformed text data in TF-IDF format for training and testing
45      ↪ respectively.
46      - vectorizer: Fitted TfidfVectorizer object.
47      """
48      #####
49      # TODO: Convert the text data to TF-IDF format and return the transformed data and the vectorizer
50      # Create a TfidfVectorizer
51      vectorizer = TfidfVectorizer()
52
53      # Fit and transform X_train
54      X_train_tfidf = vectorizer.fit_transform(X_train)
55
56      # Transform X_test
57      X_test_tfidf = vectorizer.transform(X_test)
58      #####
59      return X_train_tfidf, X_test_tfidf, vectorizer
60
61  X_train, y_train = train_ds['sentence'], train_ds['label']
62  X_test, y_test = test_ds['sentence'], test_ds['label']
63  X_train_tfidf, X_test_tfidf, vectorizer = transform_data(X_train, X_test)
64
65  """**3. Train**"""
66
67  from sklearn.model_selection import GridSearchCV
68  from sklearn.linear_model import LogisticRegression
69  from sklearn.ensemble import RandomForestClassifier
70  from sklearn.naive_bayes import MultinomialNB
71
72  def logistic_regression(X_train_tfidf, y_train):
73      """
74      Input:
75      - X_train_tfidf: Transformed text data in TF-IDF format for training.
76      - y_train: Series containing the labels for training.
77
78      Output:
79      - clf: Trained Logistic Regression model.
80      """
81      #####
82      # Define a logistic regression classifier with max_iter
```



---

```

82     clf = LogisticRegression(max_iter=1000)
83
84     # Define hyperparameters to tune
85     param_grid = {
86         'C': [0.001, 0.01, 0.1, 1, 10, 100],
87         'penalty': ['l1', 'l2']
88     }
89
90     # Perform grid search with cross-validation to find the best hyperparameters
91     grid_search = GridSearchCV(clf, param_grid, cv=5, scoring='accuracy')
92     grid_search.fit(X_train_tfidf, y_train)
93
94     # Print the best parameters
95     print("Best Parameters: ", grid_search.best_params_)
96
97     # Get the best model with the optimal hyperparameters
98     clf = grid_search.best_estimator_
99
100    # Train the final model with the entire training data
101    clf.fit(X_train_tfidf, y_train)
102
103    #####
104    return clf
105
106    def random_forest(X_train_tfidf, y_train):
107        """
108        Input:
109        - X_train_tfidf: Transformed text data in TF-IDF format for training.
110        - y_train: Series containing the labels for training.
111
112        Output:
113        - clf: Trained Random Forest classifier.
114        """
115        #####
116        # Define a Random Forest classifier
117        clf = RandomForestClassifier()
118
119        # Define hyperparameters to tune
120        param_grid = {
121            'n_estimators': [100, 200, 300], # Number of trees in the forest
122            'max_depth': [None, 10, 20], # Maximum depth of the tree
123        }
124
125        # Perform grid search with cross-validation to find the best hyperparameters
126        grid_search = GridSearchCV(clf, param_grid, cv=5, scoring='accuracy')
127        grid_search.fit(X_train_tfidf, y_train)
128
129        # Print the best parameters
130        print("Best Parameters: ", grid_search.best_params_)
131
132        # Get the best model with the optimal hyperparameters
133        clf = grid_search.best_estimator_
134
135        # Train the final model with the entire training data
136        clf.fit(X_train_tfidf, y_train)
137
138        #####
139        return clf
140
141    def naive_bayes_classifier(X_train_tfidf, y_train):
142        """
143        Input:
144        - X_train_tfidf: Transformed text data in TF-IDF format for training.
145        - y_train: Series containing the labels for training.
146
147        Output:
148        - clf: Trained Multinomial Naive Bayes classifier.
149        """
150        #####
151        # Define a Multinomial Naive Bayes classifier
152        clf = MultinomialNB()
153
154        # Define hyperparameters to tune
155        param_grid = {
156            'alpha': [0.1, 0.5, 1.0, 2.0] # Smoothing parameter (Laplace/Lidstone smoothing)
157        }
158
159        # Perform grid search with cross-validation to find the best hyperparameters
160        grid_search = GridSearchCV(clf, param_grid, cv=5, scoring='accuracy')
161        grid_search.fit(X_train_tfidf, y_train)
162
163        # Print the best parameters
164        print("Best Parameters: ", grid_search.best_params_)
165
166        # Get the best model with the optimal hyperparameters
167        clf = grid_search.best_estimator_
168

```

---

```

169     # Train the final model with the entire training data
170     clf.fit(X_train_tfidf, y_train)
171
172     #####
173     return clf
174
175     clf = logistic_regression(X_train_tfidf, y_train)
176
177     clf_rf = random_forest(X_train_tfidf, y_train)
178
179     clf_nb = naive_bayes_classifier(X_train_tfidf, y_train)
180
181     """*4. Evaluation*"""
182
183     from sklearn.metrics import accuracy_score, classification_report
184
185     def evaluate_model(clf, X_test_tfidf, y_test):
186         """
187         Input:
188         - clf: Trained Logistic Regression model.
189         - X_test_tfidf: Transformed text data in TF-IDF format for testing.
190         - y_test: Series containing the labels for testing.
191
192         Output:
193         - None (This function will print the evaluation results.)
194         """
195         #####
196         # TODO: Evaluate the model and print the results
197         # Predict the labels using the trained classifier
198         y_pred = clf.predict(X_test_tfidf)
199
200         # Calculate the accuracy
201         accuracy = accuracy_score(y_test, y_pred)
202
203         #####
204         print(f"Accuracy: {accuracy:.2f}")
205         print("Classification Report:")
206         print(classification_report(y_test, y_pred))
207
208     evaluate_model(clf, X_test_tfidf, y_test)
209
210     evaluate_model(clf_rf, X_test_tfidf, y_test)
211
212     evaluate_model(clf_nb, X_test_tfidf, y_test)

```

---

---

## Task\_2

(Biswas, 2020)

---

```
1  # -*- coding: utf-8 -*-
2  """SWE3011_41_Task2.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/12C7W-1nFm4VgbgysgY1E3tIqFN99MLtf
8
9  # SWE3011_41 Task2
10
11  **Prompt Engineering with Langchain for LLMS**
12
13  1. Based on the given code, you need to extend (modify) to a classification task using Langchain.
14  2. Since large language model baseline is required for the task, you may choose any model from OpenAI API and
   ↪ HuggingfaceHub.
15
16  **OpenAI**: Only available for paid users.
17
18  **HuggingfaceHUB**: Free to use with usage limits (reset hourly).
19
20  3. Conduct experiments and document the results in the report. Here, you should consider what kind of prompt
   ↪ design you use so please find some tutorials/resources in our homework description to obtain more
   ↪ information.
21
22  **0. Installation**
23  """
24
25  pip install datasets
26
27  pip install openai
28
29  from getpass import getpass
30  import os
31
32  # get a token: https://huggingface.co/docs/api-inference/quicktour#get-your-api-token
33  HUGGINGFACEHUB_API_TOKEN = getpass()
34
35  os.environ["HUGGINGFACEHUB_API_TOKEN"] = HUGGINGFACEHUB_API_TOKEN
36
37  pip install huggingface_hub
38
39  pip install langchain
40
41  """**1. Load Dataset**
42
43  Evaluation should be done using **provided test dataset**
44  """
45
46  from datasets import load_dataset
47  import pandas as pd
48
49  # You can use train_ds for few-shot examples
50  train_ds = load_dataset("glue", "sst2", split="train")
51
52  # Convert the dataset to a Pandas DataFrame
53  df_train = train_ds.to_pandas()
54
55  # Display the DataFrame
56  print(df_train.head())
57
58  # Evaluation should be done using test_ds
59  test_ds = load_dataset("csv", data_files="./test_dataset.csv")['train']
60
61  # Load the dataset from CSV
62  dataset_path = "/content/test_dataset.csv" # Replace with the actual path to your CSV file
63  df_test = pd.read_csv(dataset_path)
64
65  # Display the DataFrame
66  print(df_test.head())
67
68  import random
69  import numpy as np
70
71  def gen_data(df, flag=True, samples=27):
72      examples = []
73
74      n = len(df) if flag else samples
75
76      # Count occurrences of each category
77      category_counts = df["label"].value_counts()
78
79      for i in range(0, n):
```

---

```

80     # Choose category proportionally based on occurrences
81     category = np.random.choice(category_counts.index.tolist(), p=(category_counts /
82     ↪ category_counts.sum()).tolist())
83
84     # Get a random row with the chosen category
85     category_df = df[df["label"] == category]
86     random_row = category_df.sample()
87
88     obj = {
89         "messages": random_row["sentence"].iloc[0],
90         "categories": random_row["label"].iloc[0]
91     }
92     examples.append(obj)
93
94     random.shuffle(examples)
95
96     return examples
97
98     """*2. Preparing Prompt*"""
99
100    from langchain import PromptTemplate
101
102    prompt_template = '''
103    Messages: {messages}
104    Categories: {categories}
105    '''
106
107    example_prompt = PromptTemplate(
108        input_variables=["messages", "categories"],
109        template=prompt_template)
110
111    import re
112
113    def trim(s):
114        if "positive" in s.lower():
115            return 1
116        elif "negative" in s.lower():
117            return 0
118        else:
119            return -1 # or any default value for the case where no match is found
120
121    from sklearn.metrics import accuracy_score
122    from tqdm import tqdm
123    import time
124
125    def eval(extraction_chain, example):
126        preds = []
127        actual = []
128
129        for x in tqdm(example):
130            p = extraction_chain.predict(messages=x["messages"])
131            p = trim(p)
132            a = x["categories"]
133            preds.append(p)
134            actual.append(a)
135
136        accuracy = accuracy_score(actual, preds)
137        return preds, actual, accuracy
138
139    from langchain.llms import HuggingFaceHub
140    from langchain.chains import LLMChain
141
142    def initialize_llm(model_name, api_key=None):
143        """
144        Initialize the model using the langchain library.
145        """
146        llm = HuggingFaceHub(repo_id=model_name, model_kwargs={"temperature": 0.5, "max_length": 64})
147
148        return llm
149
150    """*Few-Shot*"""
151
152    from langchain import FewShotPromptTemplate
153
154    examples = gen_data(df_train, False)
155    print(examples)
156
157    # Create the FewShotPromptTemplate
158    final_prompt = FewShotPromptTemplate(
159        examples=examples,
160        example_prompt=example_prompt,
161        suffix="Messages: {messages} \nCategories: ",
162        input_variables=["messages"],
163        prefix="""You are a sentiment classifier. Classify the given text input as "positive" or "negative" based
164        ↪ on its tone.
165        Please read each sentence carefully and determine its sentiment based on common definitions.

```

---

```

165         <positive>: A positive sentiment indicates a favorable or optimistic tone. It expresses approval,
166         ↳ satisfaction, or positive emotions.
167
168         <negative>: A negative sentiment indicates an unfavorable or pessimistic tone. It expresses
169         ↳ disapproval, dissatisfaction, or negative emotions.
170
171         Please read each sentence and classify it as either "positive" or "negative" based on the provided
172         ↳ definitions.
173
174     """
175 )
176
177 # Initialize the language model
178 model_name = "google/flan-t5-xxl"
179 llm = initialize_llm(model_name)
180 val = gen_data(df_test)
181
182 # Create the LLMChain
183 extraction_chain = LLMChain(llm=llm, prompt=final_prompt, output_key="categories")
184
185 # Evaluate and print results without a new line after a comma
186 preds, actual, acc = eval(extraction_chain, val)
187 print("\nPredicted Categories:", preds)
188
189 print("Actual Categories:", actual)
190
191 print(f"Accuracy: {acc * 100:.2f}%")
192
193 from langchain import FewShotPromptTemplate
194
195 # Initialize the language model
196 model_name = "google/flan-t5-xxl"
197 llm = initialize_llm(model_name)
198
199 # List to store accuracies
200 accuracies = []
201
202 # Perform 10 iterations
203 for iteration in range(1, 4):
204     print(f"\nIteration {iteration}:")
205
206     # Generate new examples for each iteration
207     examples = gen_data(df_train, flag = False, samples=27)
208
209     # Define your FewShotPromptTemplate
210     final_prompt = FewShotPromptTemplate(
211         examples=examples,
212         example_prompt=example_prompt,
213         suffix="Messages: {messages} \nCategories: ",
214         input_variables=["messages"],
215         prefix="""You are a sentiment classifier. Classify the given text input as "positive" or "negative"
216         ↳ based on its tone.
217
218         Please read each sentence carefully and determine its sentiment based on common definitions.
219
220         <positive>: A positive sentiment indicates a favorable or optimistic tone. It expresses
221         ↳ approval, satisfaction, or positive emotions.
222
223         <negative>: A negative sentiment indicates an unfavorable or pessimistic tone. It expresses
224         ↳ disapproval, dissatisfaction, or negative emotions.
225
226         Please read each sentence and classify it as either "positive" or "negative" based on the
227         ↳ provided definitions.
228
229     """
230 )
231
232 # Create the LLMChain with the updated prompt
233 extraction_chain = LLMChain(llm=llm, prompt=final_prompt, output_key="categories")
234
235 # Evaluate and print results without a new line after a comma
236 preds, actual, acc = eval(extraction_chain, val)
237 accuracies.append(acc) # Store accuracy in the list
238 print("Predicted Categories:", preds)
239
240 print("Actual Categories:", actual)
241
242 print(f"Accuracy: {acc * 100:.2f}%")
243
244 # Print all accuracies at the end
245 print("\nAll Accuracies:", accuracies)
246
247 """**Zero-Shot**"""
248
249 from langchain import PromptTemplate
250
251 # Create the zero-shot prompt template
252 zero_shot_template = PromptTemplate(
253     input_variables=["messages"],

```

---

```

244     template="""You are a sentiment classifier. Classify the given text input as "positive" or "negative"
↪     based on its tone.
245         Please read each sentence carefully and determine its sentiment based on common definitions.
246
247         <positive>: A positive sentiment indicates a favorable or optimistic tone. It expresses
↪         approval, satisfaction, or positive emotions.
248
249         <negative>: A negative sentiment indicates an unfavorable or pessimistic tone. It expresses
↪         disapproval, dissatisfaction, or negative emotions.
250
251         Please read each sentence and classify it as either "positive" or "negative" based on the
↪         provided definitions.
252     """
253 )
254
255 # Initialize the language model
256 model_name = "google/flan-t5-xxl"
257 llm = initialize_llm(model_name)
258 val = gen_data(df_test)
259
260 # Create the LLMChain
261 extraction_chain = LLMChain(llm=llm, prompt=zero_shot_template, output_key="categories")
262
263 # Evaluate and print results
264 preds, actual, acc = eval(extraction_chain, val)
265 print("\nPredicted Categories:", preds)
266
267 print("Actual Categories:", actual)
268
269 print(f"Accuracy: {acc * 100:.2f}%", end=" ")

```

---

[illegible]

Figure 7: Result of prediction using zero-shot prompt

```
100%|██████████| 100/100 [01:07<00:00, 1.49it/s]
Predicted Categories: [1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, -1, 0, 0, 0, 0, 0, 1, 0, -1, 1,
Actual Categories: [1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0,
Accuracy: 83.00%
```

Figure 8: Result of prediction with the first execution of few-shot prompt

```
Iteration 1:
100%|██████████| 100/100 [00:28<00:00, 3.56it/s]
Predicted Categories: [0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0,
Actual Categories: [1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1,
Accuracy: 89.00%

Iteration 2:
100%|██████████| 100/100 [00:28<00:00, 3.49it/s]
Predicted Categories: [0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0,
Actual Categories: [1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1,
Accuracy: 87.00%

Iteration 3:
100%|██████████| 100/100 [00:31<00:00, 3.13it/s]
Predicted Categories: [0, 1, 0, 0, 1, 1, 1, -1, 1, -1, 0, 0, 0, 0, -1, 0, -1, 1, 0, 0, 0,
Actual Categories: [1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1,
Accuracy: 82.00%

All Accuracies: [0.89, 0.87, 0.82]
```

Figure 9: Result of prediction with subsequent executions of few-shot prompt